

NATIONAL ECONOMICS UNIVERSITY, VIETNAM



**The Power of Data Preparation:
A Data Storytelling Approach to Credit
Default Risk Prediction**

Course ID: Data Preparation and Visualization

Student ID: 11230545

Full Name: Nguyen Khanh Huyen

Student ID: 11230532

Full Name: Nguyen Thi Huong Giang

Student ID: 11230569

Full Name: Nguyen Dang Minh

Student ID: 11230585

Full Name: Nguyen Thi Ha Phuong

Student ID: 11230593

Full Name: Duong Thi Huyen Trang

Supervisor: Dr. Nguyen Tuan Long

Contents

| | |
|--|-----------|
| Part 1: Data Storytelling | 2 |
| 1 Executive Summary | 2 |
| 2 Introduction and Context | 2 |
| 2.1 The Real-World Problem | 2 |
| 2.2 The Dataset | 3 |
| 3 The Problem: Raw Data and Associated Challenges | 7 |
| 3.1 Exploring Raw Data | 7 |
| 3.2 Experiments with Raw Data | 7 |
| 4 The Solution: Data Preparation Journey | 13 |
| 5 The Impact: Clear Comparison of Results | 14 |
| 6 Call to Action and Recommendations | 14 |
| Part 2: Technical Analysis | 14 |
| 1 Storytelling Strategy Evaluation | 14 |
| 2 Principles of Data Visualization | 14 |
| 3 Technical Implementation Details | 14 |
| 4 Design Decisions and Rationale | 14 |
| 5 Lessons Learned and Best Practices | 14 |

Part 1: Data Storytelling

1 Executive Summary

In the high-stakes world of microfinance, where every lending decision affects the lives of underserved communities, the difference between success and failure often does not lie in the complexity of machine learning algorithms but in the quality of the data that feeds them. While it is tempting and common for practitioners to chase the latest sophisticated models or hyperparameter optimizations, the reality is stark: without clean, well-prepared data, even the most advanced model produces little more than garbage output, a classic case of “Garbage In, Garbage Out (GIGO).”

Microfinance institutions (MFIs) serve as financial lifelines for millions of underserved individuals, yet they face a critical dilemma: how to identify borrowers at risk of defaulting without access to traditional credit histories or comprehensive financial records. This report tells the story of how rigorous data preparation transformed a struggling credit risk prediction model into a reliable decision-making tool for MFIs.

Armed with a dataset containing borrower profiles, transaction histories, and repayment records, we initially set out to build a predictive model capable of flagging high-risk borrowers. However, our first attempts quickly revealed a universal truth familiar to any data analyst: raw data, regardless of quantity, is often more obstacle than asset.

Our analysis demonstrates that data preparation is not a minor technical step but the very foundation of reliable machine learning predictions. Through systematic understanding, cleaning, transformation, and feature engineering of microfinance transaction data, we show how proper data preparation converts mediocre models into robust decision-making tools that meaningfully reduce default risk.

The story unfolds through three critical phases. First, we expose the hidden dangers lurking in raw data, including missing values, inconsistencies, duplicate records, and noise that obscure meaningful patterns and mislead algorithms. Second, we document the transformation journey, where each data preparation technique (cleaning, transformation, reduction) progressively reveals clearer insights and stronger predictive signals. Finally, we present irrefutable evidence of impact through side-by-side comparisons that show dramatic improvements in model accuracy, reliability, and business value.

The takeaway is unambiguous: investing time and rigor in data preparation yields returns far greater than chasing complex model architectures alone. For microfinance institutions, this translates into fewer missed defaults, more responsible lending decisions, and sustainable operations that continue to serve vulnerable communities. Through data-driven storytelling, this report proves that in machine learning for credit risk assessment, data preparation is not optional; moreover, it is everything.

2 Introduction and Context

2.1 The Real-World Problem

Millions of individuals around the world face a fundamental barrier to economic opportunity: they cannot access formal credit. The reason is straightforward yet devastating that they lack traditional credit histories. Without records of past loans, credit cards, or formal banking relationships, these individuals are invisible to conventional lending

systems. Financial institutions, unable to assess their creditworthiness through standard metrics, simply reject their applications.

This credit invisibility creates a vicious cycle. Unable to secure loans from legitimate institutions, the unbanked population often turns to predatory lenders who charge exorbitant interest rates and impose exploitative terms. These untrustworthy lenders trap vulnerable borrowers in cycles of debt, undermining their financial stability, and perpetuating poverty. The consequence is not only merely individual hardship, but is also a systemic failure that locks entire communities out of economic participation.

Home Credit, a consumer finance provider committed to broadening financial inclusion, confronts this challenge head-on. Their mission is to provide positive and safe borrowing experiences for the unbanked population, ensuring that the lack of traditional credit history does not become a permanent barrier to financial opportunity.

However, the challenge is nuanced. The goal is not simply to approve more loans; it is to make the right lending decisions. Home Credit must identify clients who are capable of repayment but would be rejected by traditional systems, while simultaneously protecting vulnerable borrowers from taking on debt they cannot manage. Every false negative (rejecting a creditworthy applicant) denies someone an opportunity for economic advancement. Every false positive (approving a risky borrower) can lead to default, financial distress, and harm to both the institution and the individual.

2.2 The Dataset

The dataset provided contains comprehensive information about borrowers, their loan applications, repayment histories, and related transactional data. It includes demographic attributes, financial indicators, and behavioral information that collectively reflect client profiles. Key features include:

- **Client Information:** Gender, age, family status, education level, housing situation, income, number of children, employment details, ownership of assets such as cars or real estate.
- **Loan Details:** Loan type (cash or revolving), loan amount, annuity, goods price, application timing, and credit bureau inquiries.
- **Behavioral Indicators:** Previous payment patterns, social network influence on repayment (30 and 60 days past due), and document submissions.
- **External Sources:** Normalized scores from external data sources, regional ratings, and address consistency flags.

Each observation in the dataset corresponds to a loan application, with the target variable indicating whether the client experienced payment difficulties during the initial installments.

| Column | Description |
|--------|--------------------------|
| ID | ID of loan in our sample |

| Column | Description |
|----------------------------|---|
| TARGET | Target variable (1 - client with payment difficulties: he/she had late payment more than X days on at least one of the first Y installments of the loan in our sample, 0 - all other cases) |
| NAME_CONTRACT_TYPE | Identification if loan is cash or revolving |
| CODE_GENDER | Gender of the client |
| FLAG_OWN_CAR | Flag if the client owns a car |
| FLAG_OWN_REALTY | Flag if client owns a house or flat |
| CNT_CHILDREN | Number of children the client has |
| AMT_INCOME_TOTAL | Income of the client |
| AMT_CREDIT | Credit amount of the loan |
| AMT_ANNUITY | Loan annuity |
| AMT_GOODS_PRICE | For consumer loans it is the price of the goods for which the loan is given |
| NAME_TYPE_SUITE | Who was accompanying client when he was applying for the loan |
| NAME_INCOME_TYPE | Client's income type (businessman, working, maternity leave, ...) |
| NAME_EDUCATION_TYPE | Level of highest education the client achieved |
| NAME_FAMILY_STATUS | Family status of the client |
| NAME_HOUSING_TYPE | What is the housing situation of the client (renting, living with parents, ...) |
| REGION_POPULATION_RELATIVE | Normalized population of region where client lives (higher number means more populated region) |
| DAYS_BIRTH | Client's age in days at the time of application |
| DAYS_EMPLOYED | How many days before the application the person started current employment |
| DAYS_REGISTRATION | How many days before the application did client change his registration |
| DAYS_ID_PUBLISH | How many days before the application did client change the identity document with which he applied for the loan |
| OWN_CAR_AGE | Age of client's car |
| FLAG_MOBIL | Did client provide mobile phone (1=YES, 0=NO) |
| FLAG_EMP_PHONE | Did client provide work phone (1=YES, 0=NO) |
| FLAG_WORK_PHONE | Did client provide home phone (1=YES, 0=NO) |
| FLAG_CONT_MOBILE | Was mobile phone reachable (1=YES, 0=NO) |
| FLAG_PHONE | Did client provide home phone (1=YES, 0=NO) |
| FLAG_EMAIL | Did client provide email (1=YES, 0=NO) |

| Column | Description |
|-----------------------------|--|
| OCCUPATION_TYPE | What kind of occupation does the client have |
| CNT_FAM_MEMBERS | How many family members does client have |
| REGION_RATING_CLIENT | Our rating of the region where client lives (1,2,3) |
| REGION_RATING_CLIENT_W_CITY | Our rating of the region where client lives taking city into account (1,2,3) |
| WEEKDAY_APPR_PROCESS_START | On which day of the week did the client apply for the loan |
| HOUR_APPR_PROCESS_START | Approximately at what hour did the client apply for the loan |
| REG_REGION_NOT_LIVE_REGION | Flag if client's permanent address does not match contact address (1=different, 0=same, at region level) |
| REG_REGION_NOT_WORK_REGION | Flag if client's permanent address does not match work address (1=different, 0=same, at region level) |
| LIVE_REGION_NOT_WORK_REGION | Flag if client's contact address does not match work address (1=different, 0=same, at region level) |
| REG_CITY_NOT_LIVE_CITY | Flag if client's permanent address does not match contact address (1=different, 0=same, at city level) |
| REG_CITY_NOT_WORK_CITY | Flag if client's permanent address does not match work address (1=different, 0=same, at city level) |
| LIVE_CITY_NOT_WORK_CITY | Flag if client's contact address does not match work address (1=different, 0=same, at city level) |
| ORGANIZATION_TYPE | Type of organization where client works |
| EXT_SOURCE_1 | Normalized score from external data source |
| EXT_SOURCE_2 | Normalized score from external data source |
| EXT_SOURCE_3 | Normalized score from external data source |
| OBS_30_CNT_SOCIAL_CIRCLE | How many observation of client's social surroundings with observable 30 DPD (days past due) default |
| DEF_30_CNT_SOCIAL_CIRCLE | How many observation of client's social surroundings defaulted on 30 DPD (days past due) |
| OBS_60_CNT_SOCIAL_CIRCLE | How many observation of client's social surroundings with observable 60 DPD (days past due) default |
| DEF_60_CNT_SOCIAL_CIRCLE | How many observation of client's social surroundings defaulted on 60 DPD |
| DAYS_LAST_PHONE_CHANGE | How many days before application did client change phone |
| FLAG_DOCUMENT_2 | Did client provide document 2 |

| Column | Description |
|----------------------------|--|
| FLAG_DOCUMENT_3 | Did client provide document 3 |
| FLAG_DOCUMENT_4 | Did client provide document 4 |
| FLAG_DOCUMENT_5 | Did client provide document 5 |
| FLAG_DOCUMENT_6 | Did client provide document 6 |
| FLAG_DOCUMENT_7 | Did client provide document 7 |
| FLAG_DOCUMENT_8 | Did client provide document 8 |
| FLAG_DOCUMENT_9 | Did client provide document 9 |
| FLAG_DOCUMENT_10 | Did client provide document 10 |
| FLAG_DOCUMENT_11 | Did client provide document 11 |
| FLAG_DOCUMENT_12 | Did client provide document 12 |
| FLAG_DOCUMENT_13 | Did client provide document 13 |
| FLAG_DOCUMENT_14 | Did client provide document 14 |
| FLAG_DOCUMENT_15 | Did client provide document 15 |
| FLAG_DOCUMENT_16 | Did client provide document 16 |
| FLAG_DOCUMENT_17 | Did client provide document 17 |
| FLAG_DOCUMENT_18 | Did client provide document 18 |
| FLAG_DOCUMENT_19 | Did client provide document 19 |
| FLAG_DOCUMENT_20 | Did client provide document 20 |
| FLAG_DOCUMENT_21 | Did client provide document 21 |
| AMT_REQ_CREDIT_BUREAU_HOUR | Number of enquiries to Credit Bureau about the client one hour before application |
| AMT_REQ_CREDIT_BUREAU_DAY | Number of enquiries to Credit Bureau about the client one day before application (excluding one hour before application) |
| AMT_REQ_CREDIT_BUREAU_WEEK | Number of enquiries to Credit Bureau about the client one week before application (excluding one day before application) |
| AMT_REQ_CREDIT_BUREAU_MON | Number of enquiries to Credit Bureau about the client one month before application (excluding one week before application) |
| AMT_REQ_CREDIT_BUREAU_QRT | Number of enquiries to Credit Bureau about the client 3 months before application (excluding one month before application) |
| AMT_REQ_CREDIT_BUREAU_YEAR | Number of enquiries to Credit Bureau about the client one year before application (excluding last 3 months) |

Table 1: Description of dataset variables

3 The Problem: Raw Data and Associated Challenges

3.1 Exploring Raw Data

3.2 Experiments with Raw Data

This section evaluates the performance of three baseline machine learning models trained directly on the raw, unprocessed credit dataset. The dataset contained missing values, outliers, inconsistent feature scales, categorical variables with 50+ levels, and strong right-skewness. No cleaning, transformation, encoding, or feature engineering was applied.

Three models were tested:

- **Logistic Regression:** linear-based model
- **Random Forest:** tree-based model
- **XGBoost:** tree-based gradient boosting model

3.2.1 Logistic Regression (Linear-Based Model)

The Logistic Regression baseline demonstrates how unscaled numeric variables, high missingness, and extreme outliers disrupt linear separability. As a result, the model produces unstable and inconsistent decision boundaries.

Model Performance (Raw Data)

- Test ROC–AUC: 0.6203
- Test Accuracy: 0.6097
- Class 1 Precision: 0.0946
- Class 1 Recall: 0.5700
- Class 1 F1-score: 0.1622

Although recall for the minority class (default) appears deceptively high, precision is extremely low, indicating large numbers of false alarms.

Visual Analysis of Logistic Regression

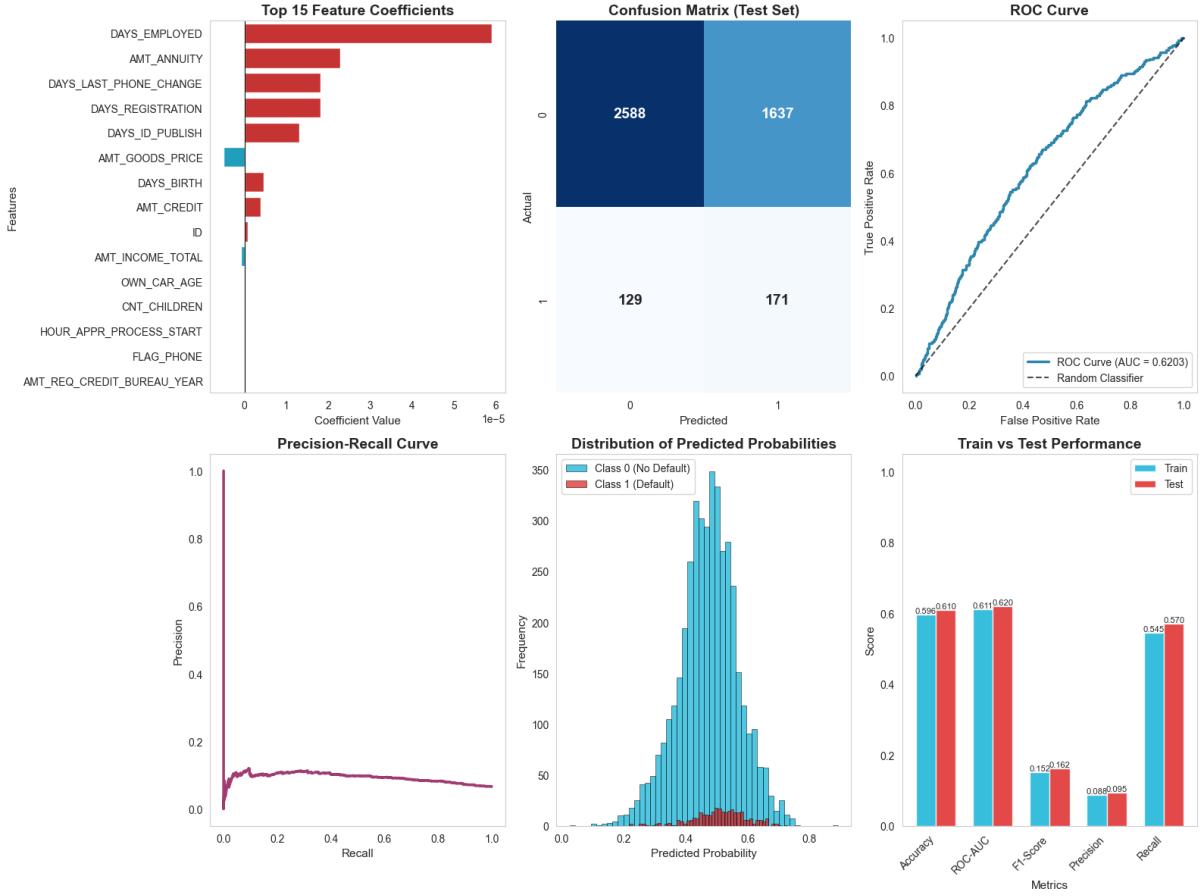


Figure 1: The visualization of Logistic Regression on Raw Data

Key Observations

- **Feature Coefficients Lack Meaning**

The coefficient plot shows extremely small magnitudes and inconsistent signs. This indicates multicollinearity, scale distortion, suppressed signal due to noise. Even domain-relevant features like EXT_SOURCE variables do not appear among top predictors.

- **Confusion Matrix: High False Positives**

The model misclassifies 1,637 non-defaulters as high-risk. This demonstrates the instability of raw linear separation.

- **ROC Curve Near the Diagonal**

The ROC curve lies close to the random classifier line, confirming low discriminative power.

- **Precision–Recall Curve Shows Almost No Predictive Value**

Precision rapidly collapses → the model cannot reliably identify default cases at any threshold.

- **Predicted Probabilities Cluster Around 0.45–0.55**

The distribution of predicted probabilities shows: lack of separation between classes, no clear decision boundary, heavy noise overshadowing signal.

Conclusion Logistic Regression confirms that linear models cannot detect meaningful patterns when raw financial data is unscaled, inconsistent, and dominated by noise.

3.2.2 Random Forest (Tree-Based Model)

Random Forest is more flexible than linear models, yet it still failed to extract meaningful relationships from the raw dataset due to noisy splits, missing values, and irrelevant high-cardinality features.

Model Performance (Raw Data)

- Test ROC-AUC: 0.7007
- Test Accuracy: 0.9193
- Class 1 Precision: 0.1905
- Class 1 Recall: 0.0667
- Class 1 F1-score: 0.0988

The model overwhelmingly predicts the majority class, yielding floor-level recall for defaulters.

Visual Analysis of Random Forest

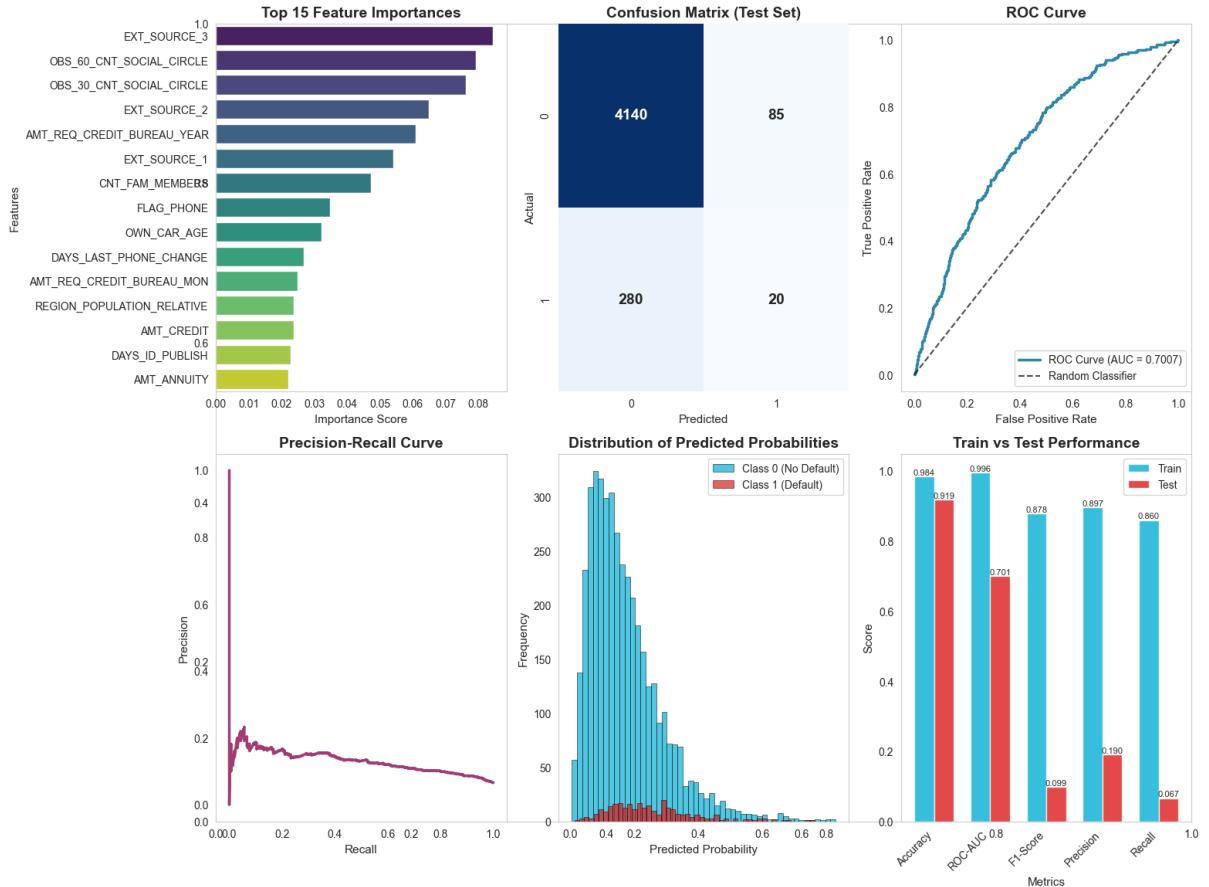


Figure 2: The visualization of Random Forest on Raw Data

Key Observations

- **Feature Importance Dominated by Noise:** Although EXT_SOURCE_3 appears as the top feature, the rest of the importance distribution is unstable and inconsistent with actual credit-risk behavior.
- **Confusion Matrix: Misses 93% of Defaults:** Only 20 out of 300 true default cases are detected. This shows the model overfits majority patterns while failing on minority detection.
- **Precision–Recall Curve Approaches Zero:** No meaningful trade-off exists; the curve indicates almost no ability to identify defaults.
- **Predicted Probabilities Extremely Skewed:** Most predictions cluster below 0.2, showing that the model saturates toward predicting non-default and that splitting rules fail due to missing values and inconsistent numeric scales.

Conclusion Random Forest is severely misled by raw noise, overfitting the majority class while missing critical minority patterns.

3.2.3 XGBoost (Tree-Based Gradient Boosting Model))

XGBoost is typically powerful at extracting nonlinear relationships. However, without imputation, encoding, or cleaning, it becomes biased toward noisy splits and missing-value default directions.

Model Performance (Raw Data)

- **Test ROC–AUC:** 0.7253
- **Test Accuracy:** 0.9335
- **Class 1 Precision:** 0.4737
- **Class 1 Recall:** 0.0300
- **Class 1 F1-score:** 0.0564

The model predicts extremely few defaulters, producing high precision but near-zero recall.

Visual Analysis of XGBoost

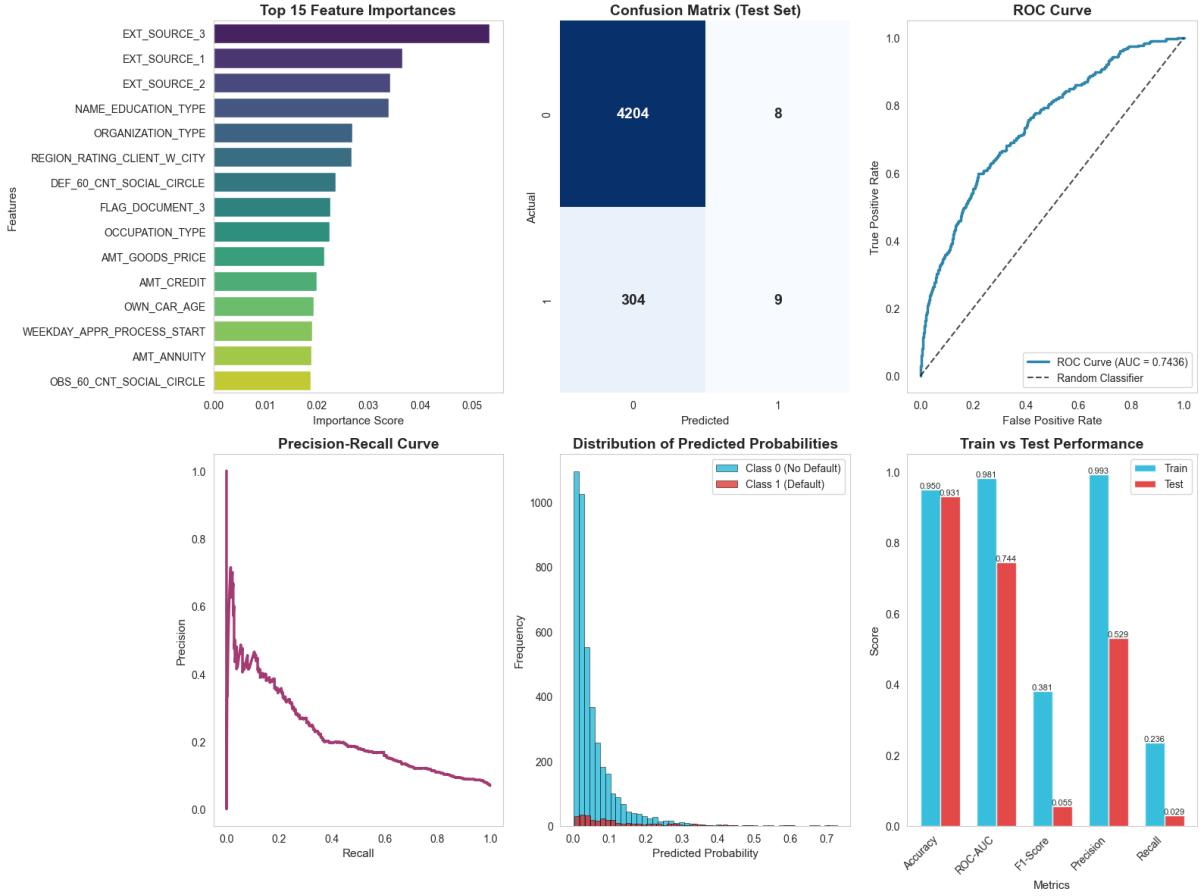


Figure 3: The visualization of on Raw Data

Key Observations

- **Feature Importance Shows Partial but Unstable Signals:** EXT_SOURCE features appear, but several high-noise variables also receive high importance, indicating brittle split logic.
- **Confusion Matrix: Almost Total Failure on Class 1:** Only 9 out of 304 default cases are correctly identified; the model collapses into predicting the majority class.
- **ROC Curve Slightly Higher Than Other Models:** Although AUC reaches 0.74, the performance does not translate into meaningful detection due to poor threshold calibration.
- **Predicted Probability Distribution Highly Skewed:** Most predicted probabilities are below 0.1, meaning the model views nearly all applicants as “very low risk” and misses critical signals.

Conclusion Even XGBoost cannot extract meaningful structure from unprocessed data—the underlying patterns are buried beneath missingness and noise.

3.2.4 Cross-Model Insights: Why Raw Data Fails

Across all three models, consistent failure patterns appear:

- **Hidden Insights:** Important financial behaviors—such as credit bureau history, late payments, and debt ratios—do not manifest in the models.
- **Visual Evidence of Noise:** Prediction probability histograms, PR curves, and ROC curves all reveal inconsistent distributions, class overlap, and missing structure.
- **Majority-Class Overlearning:** Tree models collapse into predicting nearly all observations as non-default.
- **Severe Distortion from Missing Values:** Missing EXT_SOURCE features alone cause poor splits in tree models and unstable coefficients in linear models.
- **Outliers Break Linear Relationships:** Unscaled income, loan amount, and annuity variables introduce extreme leverage effects.

The baseline results collectively show that raw data cannot support reliable credit-risk modeling: Logistic Regression misclassifies thousands of customers and fails to capture any meaningful trend. Random Forest and XGBoost severely overfit and fail to detect minority risk, despite high apparent accuracy. Visualizations clearly reveal noise-dominated distributions, unstable patterns, and poor class separation.

Therefore, data preparation is not optional. It is the prerequisite for uncovering meaningful insights and building effective ML models.

4 The Solution: Data Preparation Journey

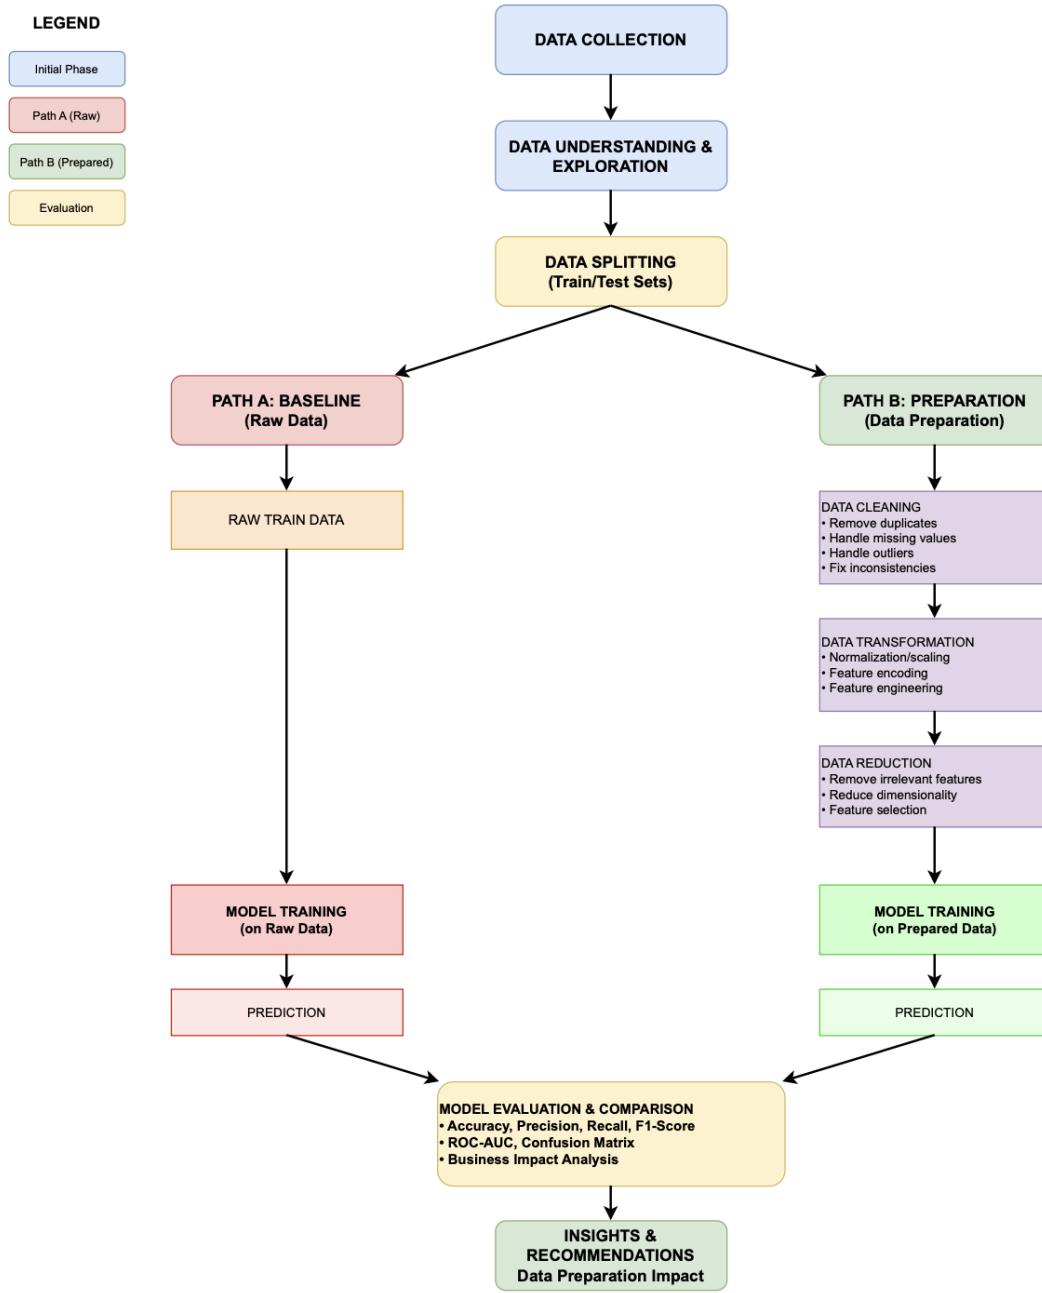


Figure 4: The flow diagram of this project

5 The Impact: Clear Comparison of Results

6 Call to Action and Recommendations

Part 2: Technical Analysis

1 Storytelling Strategy Evaluation

2 Principles of Data Visualization

3 Technical Implementation Details

4 Design Decisions and Rationale

5 Lessons Learned and Best Practices